

# Data Wrangling with *dplyr* & *tidyr*

## An introduction

Pierre-Alexandre Fonta  
R Lunchs - UniGe  
10.04.18

# Disclaimer

— — —

This presentation is a voluntary work.

My current and past institutions are therefore not responsible for its content.

# Objectives

— — —

Discover how to easily go from raw data to data ready for analytics with R.

Make your life easier with *dplyr* & *tidyr*.

Increase Data Literacy.

# Definitions

# Data Analytics

— — —

*Data Analytics is the discovery, interpretation, and communication of meaningful patterns in data.*

<https://en.wikipedia.org/wiki/Analytics>

# Data Wrangling

— — —

*Data Wrangling [...] is the process of transforming and mapping data from one 'raw' data form into another format with the intent of making it more appropriate and valuable for a variety of downstream purposes such as analytics.*

[https://en.wikipedia.org/wiki/Data\\_wrangling](https://en.wikipedia.org/wiki/Data_wrangling)

# Data Literacy

— — —

*Data Literacy is the ability to read, understand, create and communicate data as information. [...] [D]ata literacy focuses on the competencies involved in working with data.*

*As data collection and sharing become routine and data analysis [...] become[s] common ideas in the news, business, government and society, it becomes more and more important for students, citizens, and readers to have some data literacy.*

<https://en.wikipedia.org/wiki/Analytics>

**From the field**



# Personal experience - Various data

— — —

## numbers + texts + networks

- bibliographic databases (Web of Science, PubMed)

## numbers

- serious game (RTS Tabula Rasa)
- economic data (IHS Life Sciences)
- operations data (company)
- statistical data (INSEE)

## numbers + networks

- social networks (Twitter)

## states / events sequences

- biographical longitudinal data (TraMineR)

## numbers + texts


- semantic text annotations

## texts

- press releases
- web pages / blogs
- CVs (internal)
- job boards (Indeed)

# Personal experience - Various environments

— — —



# Tips - The Zen of Python

-- --

*Beautiful is better than ugly.*

*Explicit is better than implicit.*

*Simple is better than complex.*

*Flat is better than nested.*

*Readability counts.*

*Special cases aren't special enough to break the rules.*

*Although practicality beats purity.*

*In the face of ambiguity, refuse the temptation to guess.*

*There should be one - and preferably only one - obvious way to do it.*

*If the implementation is hard to explain, it's a bad idea.*

*If the implementation is easy to explain, it may be a good idea.*

<https://www.python.org/dev/peps/pep-0020/#id3>

# Tips - Might be obvious, but...

1/2

— — —

We need to understand the logic tomorrow. Our colleagues too.

Reusing variables part of the transformation logic is almost 'irresponsible'.

Modifying a logic used several times but not generalized as a function is a nightmare.

The dataset structure and the data should not be modified by the same logic / code.

# Tips - Might be obvious, but...

2/2

— — —

The time spend to format the logic to present it to non expert people is significant.

No time for implementing or optimizing the operations at the dataset level.

Data rule. We need interactivity to see them before and after.

We need a 'notebook' to capture methodology and results.

*tibble*

A modern *data.frame*!

# At the beginning of the tidyverse were...

— — —

## Tibbles!

Comparing to `data.frames`, they:

- complain when a variable does not exist
- don't change variable names or types
- don't do partial matching
- have a better `print()`

*tidyr*  
Create tidy data!



# Tidy data

— — —

Each variable is in a **column**.

Each observation is a **row**.

Each value is a **cell**.

# *tidyr* - go from wide to long


— — —

**gather**(data, key, value, ..., na.rm = FALSE,  
convert = FALSE, factor\_key = FALSE)

gather() moves column names into a **key** column, gathering the column values into a single **value** column.

table4a

country	1999	2000
A	0.7K	2K
B	37K	80K
C	212K	213K



country	year	cases
A	1999	0.7K
B	1999	37K
C	1999	212K
A	2000	2K
B	2000	80K
C	2000	213K

key value

```
gather(table4a, `1999`, `2000`,  
key = "year", value = "cases")
```

# *tidyr* - go from long to wide

— — —

**spread**(data, key, value, fill = NA, convert = FALSE,  
drop = TRUE, sep = NULL)

spread() moves the unique values of a **key**  
column into the column names, spreading the  
values of a **value** column across the new columns.

table2

country	year	type	count
A	1999	cases	0.7K
A	1999	pop	19M
A	2000	cases	2K
A	2000	pop	20M
B	1999	cases	37K
B	1999	pop	172M
B	2000	cases	80K
B	2000	pop	174M
C	1999	cases	212K
C	1999	pop	1T
C	2000	cases	213K
C	2000	pop	1T

key value

spread(table2, type, count)

country	year	cases	pop
A	1999	0.7K	19M
A	2000	2K	20M
B	1999	37K	172M
B	2000	80K	174M
C	1999	212K	1T
C	2000	213K	1T

# *tidyr* - split a column into columns

— — —

**separate**(data, col, into, sep = "[^:alnum:]]+", remove = TRUE, convert = FALSE, extra = "warn", fill = "warn", ...)

Separate each cell in a column to make several columns.

table3

country	year	rate	country	year	cases	pop
A	1999	0.7K/19M	A	1999	0.7K	19M
A	2000	2K/20M	A	2000	2K	20M
B	1999	37K/172M	B	1999	37K	172
B	2000	80K/174M	B	2000	80K	174
C	1999	212K/1T	C	1999	212K	1T
C	2000	213K/1T	C	2000	213K	1T

*separate(table3, rate,  
into = c("cases", "pop"))*

# *tidyr* - handle missing values

— — —

## **drop\_na(data, ...)**

Drop rows containing  
NA's in ... columns.

X

x1	x2
A	1
B	NA
C	NA
D	3
E	NA

 → 

x1	x2
A	1
D	3

*drop\_na(x, x2)*

## **fill(data, ..., .direction = c("down", "up"))**

Fill in NA's in ... columns with most  
recent non-NA values.

X

x1	x2
A	1
B	NA
C	NA
D	3
E	NA

 → 

x1	x2
A	1
B	1
C	1
D	3
E	3

*fill(x, x2)*

## **replace\_na(data,**

**replace = list(), ...)**

Replace NA's by column.

X

x1	x2
A	1
B	NA
C	NA
D	3
E	NA

 → 

x1	x2
A	1
B	2
C	2
D	3
E	2

*replace\_na(x, list(x2 = 2))*

*dplyr*  
Handle your data!

# Pipes

---

$x \%>\% f(y)$  is equivalent to  $f(x, y)$

# *dplyr* - summarise rows

— — —



**summarise**(.data, ...)  
Compute table of summaries.  
*summarise(mtcars, avg = mean(mpg))*



**count**(x, ..., wt = NULL, sort = FALSE)  
Count number of rows in each group defined by the variables in ... Also **tally**().  
*count(iris, Species)*

## VARIATIONS

**summarise\_all()** - Apply funs to every column.

**summarise\_at()** - Apply funs to specific columns.

**summarise\_if()** - Apply funs to all cols of one type.



# *dplyr* - group rows by variable(s)

— — —



```
mtcars %>%  
  group_by(cyl) %>%  
  summarise(avg = mean(mpg))
```

**group\_by**(.data, ..., add = FALSE)

Returns copy of table  
grouped by ...

```
g_iris <- group_by(iris, Species)
```


**ungroup**(x, ...)

# *dplyr* - keep specific rows

— — —



**filter**(.data, ...) Extract rows that meet logical criteria. *filter(iris, Sepal.Length > 7)*



**distinct**(.data, ..., .keep\_all = FALSE) Remove rows with duplicate values. *distinct(iris, Species)*



**slice**(.data, ...) Select rows by position. *slice(iris, 10:15)*

**top\_n**(x, n, wt) Select and order top n entries (by group if grouped data). *top\_n(iris, 5, Sepal.Width)*

# *dplyr* - order rows

— — —



**arrange(.data, ...)** Order rows by values of a column or columns (low to high), use with **desc()** to order from high to low.  
`arrange(mtcars, mpg)`  
`arrange(mtcars, desc(mpg))`

# *dplyr* - keep specific variables

— — —



**pull(.data, var = -1)** Extract column values as a vector. Choose by name or index.  
*pull(iris, Sepal.Length)*



**select(.data, ...)**  
Extract columns as a table. Also **select\_if()**.  
*select(iris, Sepal.Length, Species)*

**Use these helpers with select (),**  
*e.g. select(iris, starts\_with("Sepal"))*

**contains(match)**

**ends\_with(match)**

**matches(match)      starts\_with(match)**

# *dplyr* - rename variable

— — —



**rename**(.data, ...) Rename columns.  
*rename(iris, Length = Sepal.Length)*

# *dplyr* - add new variable(s)

— — —




**mutate(.data, ...)**  
Compute new column(s).  
*mutate(mtcars, gpm = 1/mpg)*



**mutate\_all(.tbl, .funs, ...)** Apply funs to every column. Use with **funs()**. Also **mutate\_if()**.  
*mutate\_all(faithful, funs(log(.), log2(.)))*  
*mutate\_if(iris, is.numeric, funs(log(.)))*



**mutate\_at(.tbl, .cols, .funs, ...)** Apply funs to specific columns. Use with **funs()**, **vars()** and the helper functions for **select()**.  
*mutate\_at(iris, vars(-Species), funs(log(.)))*



**add\_column(.data, ..., .before = NULL, .after = NULL)** Add new column(s). Also **add\_count()**, **add\_tally()**. *add\_column(mtcars, new = 1:32)*

# dplyr - join datasets

— — —

A	B	C	D
a	t	1	3
b	u	2	2
c	v	3	NA

**left\_join**(x, y, by = NULL, copy=FALSE, suffix=c(".x",".y"),...) Join matching values from y to x.

A	B	C	D
a	t	1	3
b	u	2	2
d	w	NA	1

**right\_join**(x, y, by = NULL, copy = FALSE, suffix=c(".x",".y"),...) Join matching values from x to y.

A	B	C	D
a	t	1	3
b	u	2	2

**inner\_join**(x, y, by = NULL, copy = FALSE, suffix=c(".x",".y"),...) Join data. Retain only rows with matches.

A	B	C	D
a	t	1	3
b	u	2	2
c	v	3	NA
d	w	NA	1

**full\_join**(x, y, by = NULL, copy=FALSE, suffix=c(".x",".y"),...) Join data. Retain all values, all rows.

A	B.x	C	B.y	D
a	t	1	t	3
b	u	2	u	2
c	v	3	NA	NA

Use **by = c("col1", "col2")** to specify the column(s) to match on.  
`left_join(x, y, by = "A")`

A.x	B.x	C	A.y	B.y
a	t	1	d	w
b	u	2	b	u
c	v	3	a	t

Use a named vector, **by = c("col1" = "col2")**, to match on columns with different names in each data set.  
`left_join(x, y, by = c("C" = "D"))`

A1	B1	C	A2	B2
a	t	1	d	w
b	u	2	b	u
c	v	3	a	t

Use **suffix** to specify suffix to give to duplicate column names.  
`left_join(x, y, by = c("C" = "D"), suffix = c("1", "2"))`

# *dplyr* - set operations on rows

— — —

A	B	C
c	v	3

## **intersect(x, y, ...)**

Rows that appear in both x and y.



A	B	C
a	t	1
b	u	2

## **setdiff(x, y, ...)**

Rows that appear in x but not y.



A	B	C
a	t	1
b	u	2
c	v	3
d	w	4

## **union(x, y, ...)**

Rows that appear in x or y.  
(Duplicates removed). `union_all()`  
retains duplicates.



A	B	C
a	t	1
b	u	2

## **semi\_join(x, y, by = NULL, ...)**

Return rows of x that have a match in y.  
USEFUL TO SEE WHAT WILL BE JOINED.

A	B	C
c	v	3

## **anti\_join(x, y, by = NULL, ...)**

Return rows of x that do not have a  
match in y. USEFUL TO SEE WHAT WILL  
NOT BE JOINED.

Use **setequal()** to test whether two data sets  
contain the exact same rows (in any order).



# *dplyr* - functions for summarise()

— — —

## COUNTS

**dplyr::n()** - number of values/rows

**dplyr::n\_distinct()** - # of uniques

**sum(!is.na())** - # of non-NA's

## LOCATION

**mean()** - mean, also **mean(!is.na())**

**median()** - median

## LOGICALS

**mean()** - Proportion of TRUE's

**sum()** - # of TRUE's

## POSITION/ORDER

**dplyr::first()** - first value

**dplyr::last()** - last value

**dplyr::nth()** - value in nth location of vector

## RANK

**quantile()** - nth quantile

**min()** - minimum value

**max()** - maximum value

## SPREAD

**IQR()** - Inter-Quartile Range

**mad()** - median absolute deviation

**sd()** - standard deviation

**var()** - variance

# *dplyr* - functions for mutate()

— — —

## OFFSETS

`dplyr::lag()` - Offset elements by 1  
`dplyr::lead()` - Offset elements by -1

## CUMULATIVE AGGREGATES

`dplyr::cumall()` - Cumulative all()  
`dplyr::cumany()` - Cumulative any()  
    `cummax()` - Cumulative max()  
`dplyr::cummean()` - Cumulative mean()  
    `cummin()` - Cumulative min()  
    `cumprod()` - Cumulative prod()  
    `cumsum()` - Cumulative sum()

## RANKINGS

`dplyr::cume_dist()` - Proportion of all values <=

`dplyr::dense_rank()` - rank with ties = min, no gaps

`dplyr::min_rank()` - rank with ties = min

`dplyr::ntile()` - bins into n bins

`dplyr::percent_rank()` - min\_rank scaled to [0,1]

`dplyr::row_number()` - rank with ties = "first"

## MATH

+, -, \*, /, ^, %/%, %% - arithmetic ops  
`log()`, `log2()`, `log10()` - logs  
<, <=, >, >=, !=, == - logical comparisons

`dplyr::between()` - x >= left & x <= right

`dplyr::near()` - safe == for floating point numbers

## MISC

`dplyr::case_when()` - multi-case if\_else()

`dplyr::coalesce()` - first non-NA values by element across a set of vectors

`dplyr::if_else()` - element-wise if() + else()

`dplyr::na_if()` - replace specific values with NA

`pmax()` - element-wise max()

`pmin()` - element-wise min()

`dplyr::recode()` - Vectorized switch()

`dplyr::recode_factor()` - Vectorized switch() for factors

*Jupyter*

A multilanguage notebook!

# Jupyter Notebook

— — —

<https://jupyter.org>

jupyter R Lunchs -- dplyr & tidyr -- 10-04-18 (unsaved changes) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted R

Challenge data.table!

```
In [126]: tbl1 = tibble(x = c(rep(1, 4), rep(2, 4)),  
                     y = c("a", "a", "b", "b", "a", "a", "b", "b"),  
                     z = 1:8)
```

```
In [127]: tbl2 = tibble(x = 1:2, y = c("a", "b"), mul = 4:3)
```

```
In [128]: tbl1
```

x	y	z
1	a	1
1	a	2
1	b	3
1	b	4
2	a	5
2	a	6
2	b	7
2	b	8

```
In [129]: tbl2
```

x	y	mul
1	a	4
2	b	3

```
In [140]: tbl1 %>%  
  group_by(x, y) %>%  
  summarise(sum_z = sum(z)) %>%  
  inner_join(tbl2, by = c("x", "y")) %>%  
  mutate(sum_mul = sum_z * mul) %>%  
  select(x, y, sum_mul)
```

x	y	sum_mul
1	a	12
2	b	45

# JupyterLab - Coming Soon (beta)

File Edit View Run Kernel Tabs Settings Help

Files

- + notebooks
- Name Last Modified
- Data.ipynb an hour ago
- Fasta.ipynb a day ago
- Julia.ipynb a day ago
- Lorenz.ipynb seconds ago**
- R.ipynb a day ago
- iris.csv a day ago
- lightning.json 9 days ago
- lorenz.py 3 minutes ago

Running

Commands

Cell Tools

Output View

lorenz.py

Python 3

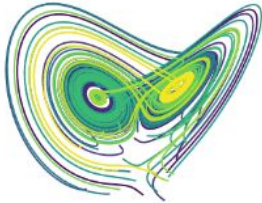
In this Notebook we explore the Lorenz system of differential equations:

$$\begin{aligned}\dot{x} &= \sigma(y - x) \\ \dot{y} &= \rho x - y - xz \\ \dot{z} &= -\beta z + xy\end{aligned}$$

Let's call the function once to view the solutions. For this set of parameters, we see the trajectories swirling around two points, called attractors.

In [4]: `from lorenz import solve_lorenz  
t, x_t = solve_lorenz(N=10)`

sigma 10.00  
beta 2.67  
rho 28.00



```
9 def solve_lorenz(N=10, max_time=4.0, sigma=10.0, beta=8./3, rho=28.0):  
10     """Plot a solution to the Lorenz differential equations."""  
11     fig = plt.figure()  
12     ax = fig.add_axes([0, 0, 1, 1], projection='3d')  
13     ax.axis('off')  
14  
15     # prepare the axes limits  
16     ax.set_xlim((-25, 25))  
17     ax.set_ylim((-35, 35))  
18     ax.set_zlim((5, 55))  
19  
20     def lorenz_deriv(x_y_z, t0, sigma=sigma, beta=beta, rho=rho):  
21         """Compute the time-derivative of a Lorenz system."""  
22         x, y, z = x_y_z  
23         return [sigma * (y - x), x * (rho - z) - y, x * y - beta * z]  
24  
25     # Choose random starting points, uniformly distributed from -15 to 15  
26     np.random.seed(1)  
27     x0 = -15 + 30 * np.random.random((N, 3))  
28
```

<https://jupyterlab.readthedocs.io/en/latest/images/jupyterlab.png>

Demo

# Takeaways

# Takeaways

— — —

## *dplyr* / *tidyr* syntax

- beautiful (~~ugly~~)
- explicit (~~implicit~~)
- simple (~~complex~~)
- readable, easy to explain

## tidy data

- flat (~~nested~~)
- analytics oriented

## Jupyter

- interactive
- code, outputs, and comments in one place



# To go further

— — —

## Databases (*dbplyr*)

<http://dbplyr.tidyverse.org>

## Programming

<http://dplyr.tidyverse.org/articles/programming.html>

## Windows functions

<http://dplyr.tidyverse.org/articles/window-functions.html>

# References

— — —

***tibble***

<http://tibble.tidyverse.org>

***dplyr***

<http://dplyr.tidyverse.org>

<https://github.com/rstudio/cheatsheets/raw/master/data-transformation.pdf>

***tidyr***

<http://tidyr.tidyverse.org>

<https://github.com/rstudio/cheatsheets/raw/master/data-import.pdf>